

Structured Query Verbalization with Large Language Models



SCAN ME

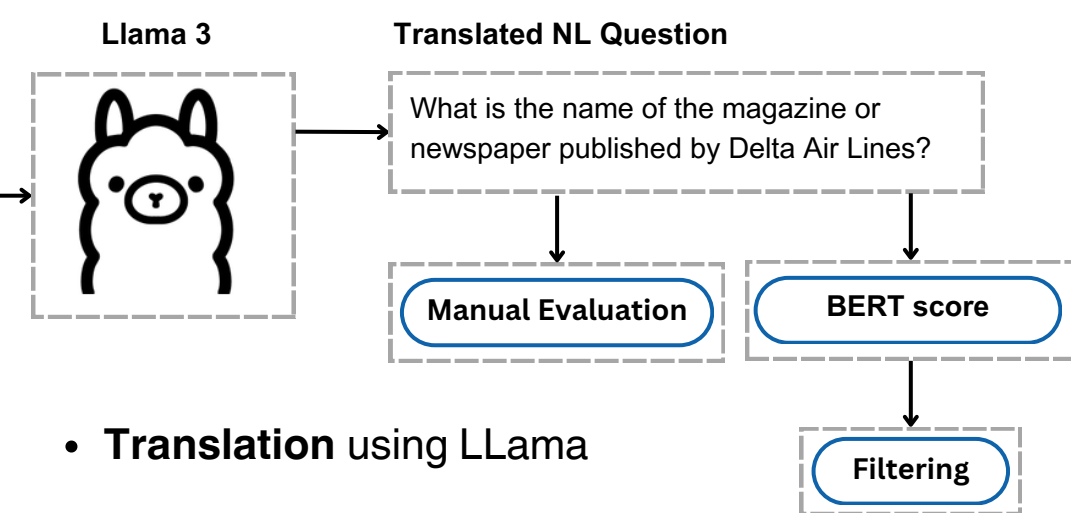
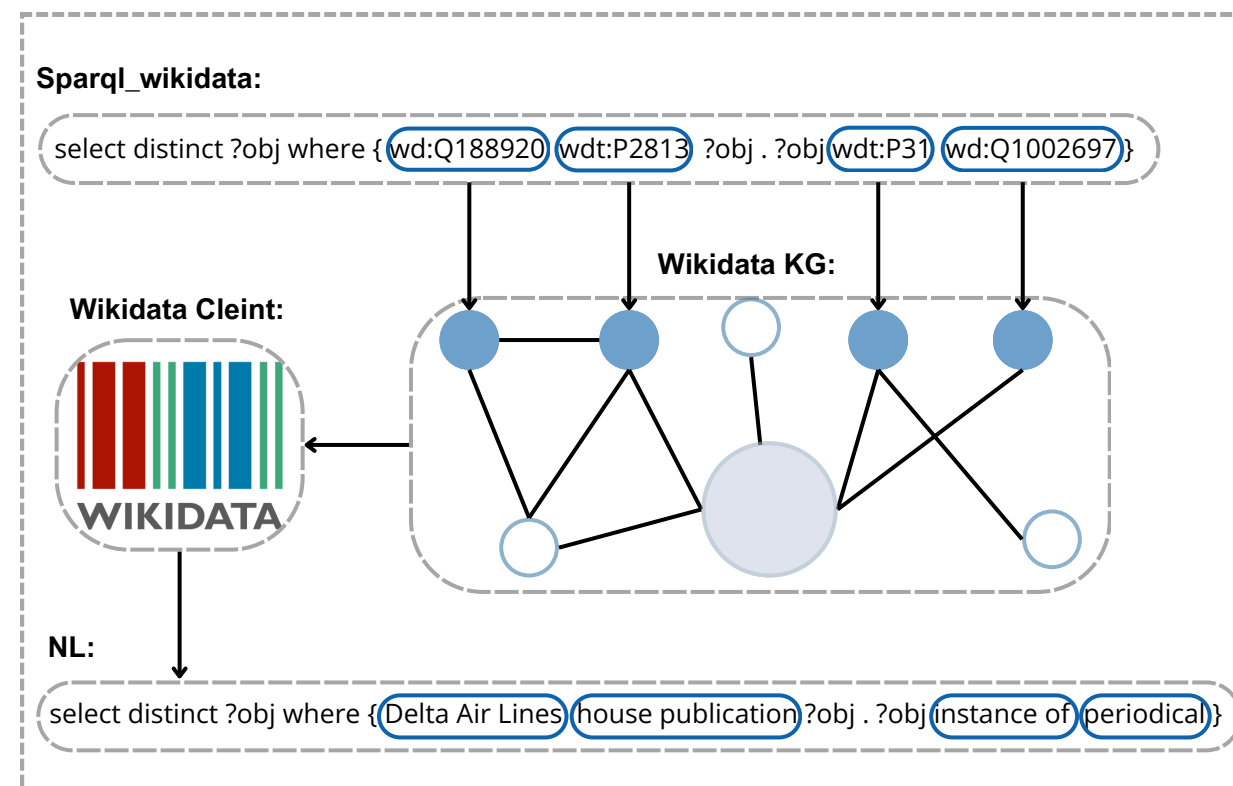
This project aims to provide a way of generating data for training LLMs for the specific use case of creating SPARQL queries from natural language.

Authors: Patrik Valach, Louisa Siebel

Advisors: M.Sc Tim Schwabe

Approach

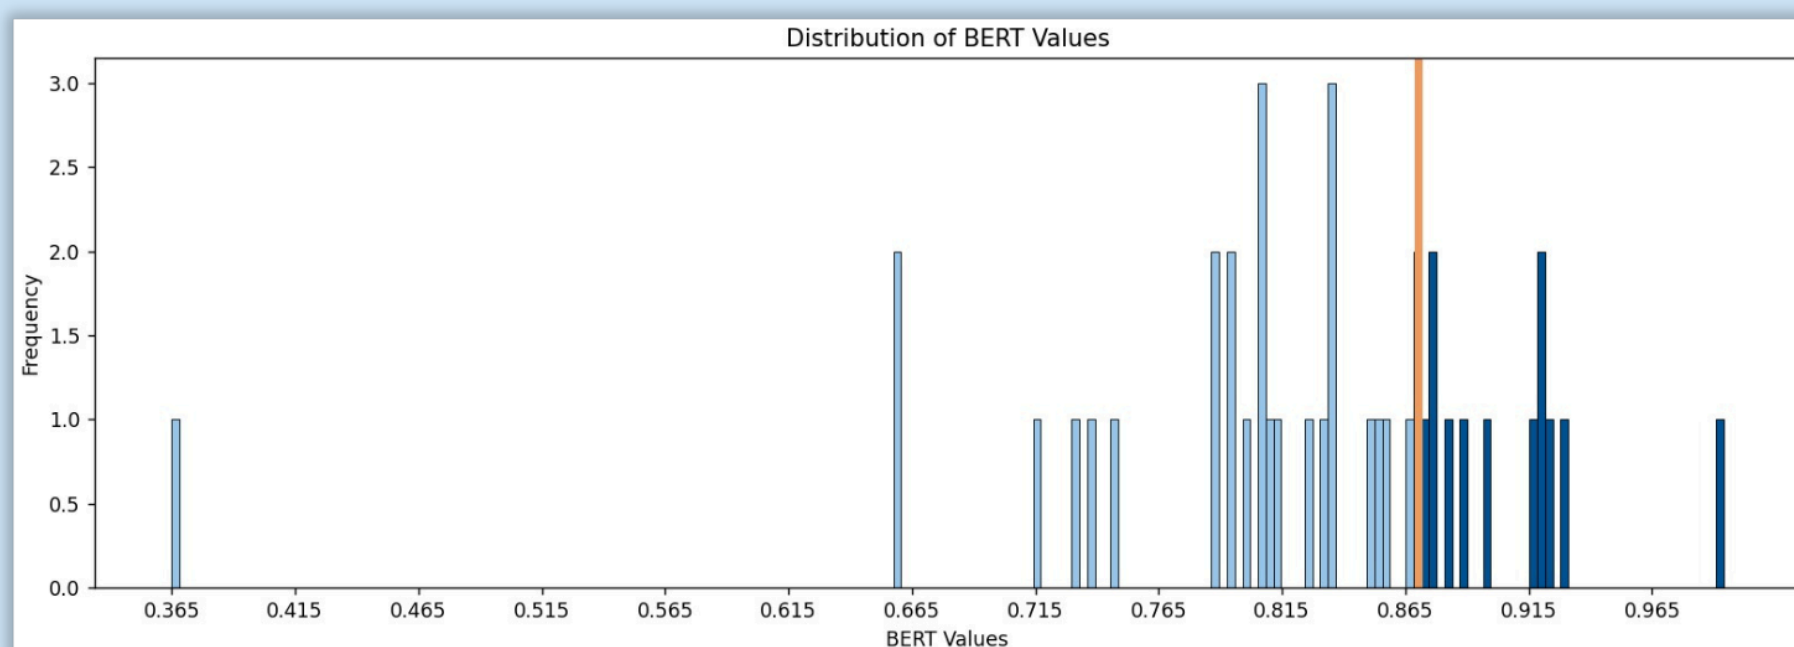
- Generating a **natural language question** from a **SPARQL query**
- Using the lc_Quad NLP-Query Database
- Map **ambiguous entities and properties** to their **natural language** form through the Wikidata client




- Translation using LLama
- Evaluation of translation performance :
 - Manual
 - BERT score
- Filtering for True Positives

Results: BERT

- BERT has proven to be the most **reliable metric** in determining the **semantic equivalence**

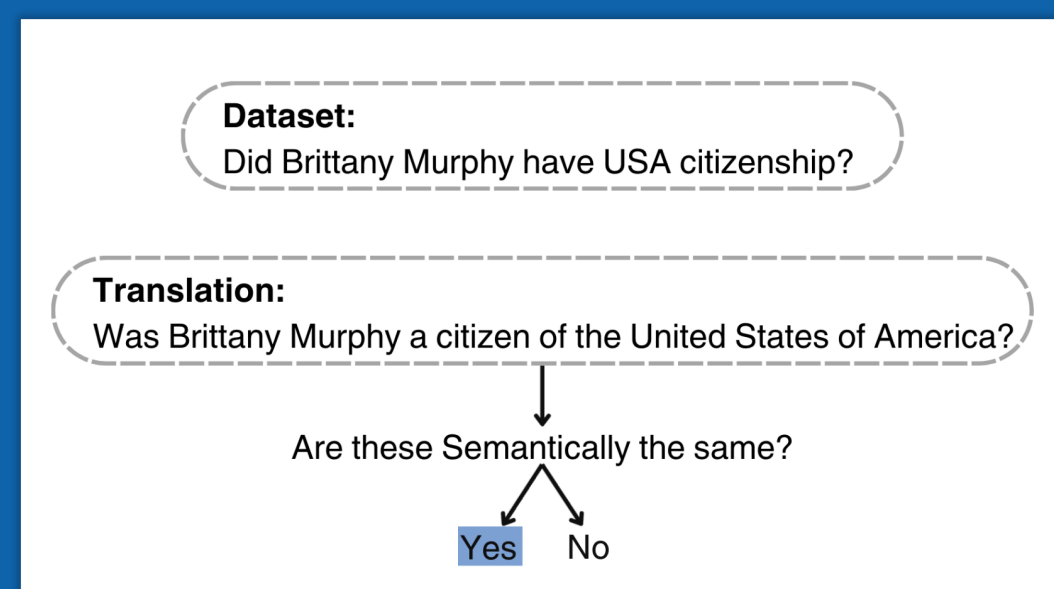


- Minimum value where we get only **True Positives** is **0,87**
 - When **filtering** with a **BERT score > 0.87** we obtain **only correct** translations




- Filter creates a set of **only correct** translations
- Filtering result:** $\approx 1/3$ of the original dataset size.
- Some **correct translations** are **sacrificed** for a truly **perfect result**

Results: Manual Inspection



- Result from manual inspection is either **0** or **1**
- Helped **establish a metric** for finding the right **filter**



- Used to **evaluate** the models' **performance**
 - 76%** of translations are **correct**
- Used to **set the threshold** for BERT

Conclusion

- Translation task is very successful: **76%**,
- We have established **metric** by which we can filter out only **true positives** if we have the ground truth
- Altogether, the project successfully **fulfills** the **goal** and can be used to **create new dataset**

Future Work

- Develop a method to **automatically** specify if a translation is **correct** without relying on the **ground truth**
 - E.g. in a **production environment** where translations need to be **generated at scale**